

Shokumu keirekisho (職務経歴書): detailed career history

Name: S Prajwall Narayana

Career Summary (職務要約)

Applicable Experience and Technical Skills (活かせる経験・知識・技術)

AI and LLM systems	Agentic systems, MCP orchestration, multi-agent workflows, RAG, LangChain, LangGraph, Claude API, LLM integration, evaluation harnesses
Infrastructure and DevOps	CI/CD, OpenTofu (IaC), Ansible, Docker, Linux, Linode, observability and alerting, incident runbooks, Zitadel (SSO/IAM), RBAC, Git
Languages and backend	Python, FastAPI, REST APIs, PostgreSQL, SQL, Swift, Rust, shell scripting
ML and computer vision	PyTorch, TensorFlow, YOLOv8, OpenCV, CNN, ResNet, quantized model deployment

Colligence Research

November 2025 – Present · Bengaluru, India

Business of the company	TODO: one line describing what the company does (a Japanese career history conventionally carries a short company outline)
Organisation size and position	Founding team; four people at the time the deployment work described below was done.
Type of employment	TODO: type of employment (permanent / contract, etc.)

AI Engineer

February 2026 – Present

- Work across both halves of the product: the AI and product engineering described below, together with end-to-end ownership of the platform it runs on.
- Established the DevOps/SRE function from scratch and now lead it: CI/CD, infrastructure-as-code with OpenTofu and Ansible on Linode, observability, alerting and incident runbooks.
- Replaced the manual GitHub deploy workflows the team had been using, which reduced deployment errors for the four-person founding team.
- Operate the production stack end to end: Docker, Zitadel authentication, and monitored deployments with self-healing alerting.
- Own system architecture across infrastructure and tooling.
- Building a retrieval-augmented system into the infrastructure layer so that infrastructure-as-code and SRE work (runbook lookup, change context, incident response) is grounded in the organisation's own configuration and history. In progress.

Associate AI Engineer

December 2025 – February 2026

- Led a small team building the company's first multi-agent LLM pipeline for ERP automation over MCP, keeping long-running, multi-step workflows on track without human intervention.
- Built a code-generation service that parses structured input specifications and emits deployment-ready web applications.

Associate AI Engineer Intern

November 2025 – December 2025

- Prototyped the first LLM integrations and agent workflows, which the production multi-agent pipeline was later built on.

Technologies: Python, FastAPI, MCP, LangChain, LangGraph, Claude API, multi-agent orchestration, OpenTofu, Ansible, Docker, Linode, Linux, Zitadel, CI/CD

RVCE Centre of Excellence

October 2024 – December 2024 · Remote

Research Intern, AI and Deep Learning

- Built a deepfake-detection model from scratch for manipulated-face video, using no external detection APIs.
- Trained and evaluated the model on the FaceForensics++ and DFDC datasets.

Technologies: Python, PyTorch, TensorFlow, CNN, computer vision, FaceForensics++, DFDC

Ernst & Young India

August 2024 – September 2024 · Bengaluru, India

Associate Consultant Intern, Enterprise Cybersecurity

- Shadowed senior consultants on cloud-security and identity and access management work.
- Gained exposure to the TOGAF and IFTAS frameworks.

Exposure: Cloud security, IAM, TOGAF, IFTAS

Selected Projects (主な担当プロジェクト)

Arogya-Sathi: offline-capable multilingual AI healthcare platform

January 2025 – June 2025

Symptom checking, X-ray fracture detection, prescription summarisation and a wellness chatbot, running fully offline on quantized models and demonstrated on a Raspberry Pi. A deterministic, rule-based drug-interaction and crisis-detection safety layer sits underneath the model and the LLM cannot override it. A 62-case labelled evaluation harness scores those safety layers and runs as a CI merge gate at recall 1.00; building it surfaced a real missed self-harm disclosure. A LangGraph pipeline turns an uploaded medical report into a plain-language explanation in the patient's own language with a self-verification loop, and requests are routed across local Ollama models. Third place, college final-year showcase. github.com/Developer1010x/Arogya-Sathi

Measured: crisis detection recall 1.00 / precision 0.93 · emergency 1.00 / 1.00 · drug screening 1.00 / 1.00

Technologies: LangGraph, Ollama, quantized models, evaluation harness, CI gate, Raspberry Pi

agentic-devops: AI-powered CI/CD pipeline generator and monitor

Generates and monitors CI/CD pipelines for GitHub and GitLab repositories. github.com/Developer1010x/agentic-devops

Knotes Central: student resource platform

A student resource platform, adopted across 12 or more departments. github.com/Developer1010x/KnotesCentral-Source-Code

pybridge.ai: AI access from a phone

Routes AI access across WhatsApp, Telegram and email so it can be driven from a phone. github.com/Developer1010x/pybridge.ai

LLM Terminal: terminal-based LLM assistant

A lightweight terminal-based assistant for working with LLMs from the command line. github.com/Developer1010x/LLM_Terminal

Audio Deepfake Detection: synthetic speech detection

Flags AI-synthesised or cloned speech. github.com/Developer1010x/Deepfake-Audio-and-AI-Content-Detection

Publications (研究業績)

EduConnect: A Smart Solution for Streamlined School Administrations.

Springer ICTIS 2025, Bangkok (ICT for Intelligent Systems, Vol. 9).

An Empirical Study of ResNet50 Hyperparameter Tuning for Plant Disease Classification.

IEEE Xplore, November 2024.
ieeexplore.ieee.org/document/10816835

Web-Server Controlled Rover with Robotic Arm and Object Detection.

IEEE Xplore, November 2024.
ieeexplore.ieee.org/document/10816816

Two further papers are under review.

Education (学歴)

December 2021 – August 2025	R V College of Engineering (Visvesvaraya Technological University), Bengaluru, India: B.E. in Computer Science and Engineering
-----------------------------	--

Languages and Qualifications (語学・資格)

Japanese	TODO: confirm Japanese-language ability with the applicant before entering anything here (e.g. studying / JLPT N- / none). Do not enter anything that cannot be confirmed. Japanese ability is screened on for nearly every role in Japan, including English-speaking ones. Do not enter a level that has not been awarded.
English	TODO: confirm with the applicant how he wishes to state his English. Keep it to documented use at work; do not print a test score he has not sat.
Certifications	TODO: qualifications held. If there are none, write "None".

Work Eligibility in Japan (就労資格)

Currently resident in Bengaluru, India, with no existing right to work in Japan. The applicable status of residence is **Engineer / Specialist in Humanities / International Services** (技術・人文知識・国際業務), the standard category for software engineers hired from overseas; a B.E. in Computer Science is a qualifying degree for it. The Certificate of Eligibility (COE) is applied for by the sponsoring employer. Available to relocate to Japan.

Self-introduction (自己PR)

What I would bring is the combination rather than either half of it: I build the agent and LLM systems and I also run the platform they are deployed on, so the person who ships a workflow is the same person who is responsible for it in production. At Colligence Research that has meant setting up the delivery and monitoring practices a founding team did not yet have, while continuing to deliver product features alongside them. I try to make correctness checkable rather than assumed: in Arogya-Sathi the safety behaviour is enforced by a deterministic layer the model cannot override and verified by a labelled evaluation set that blocks a merge if it regresses. I still have a great deal to learn, and I would welcome the opportunity to do so within your team.

End of document (以上)