

職務経歴書

Shokumu keirekisho — detailed career history

TODO: 提出年月日 年 月 日現在

氏名: S Prajwall Narayana

インド・バンガロール在住 (Bengaluru, India) | 日本への転居可 (open to relocation to Japan) | prajwallnarayana@gmail.com | +91 7760604439 | github.com/Developer1010x | sprajwallnarayana.vercel.app

職務要約 Career summary

AI Engineer on the founding team at Colligence Research in Bengaluru. My work covers both halves of the system. On the product side I build agentic and LLM pipelines, a code-generation service that turns structured input specifications into deployment-ready web applications, and the backend APIs behind them. On the platform side I established the DevOps/SRE function at the company from scratch and now lead it: CI/CD, infrastructure-as-code with OpenTofu and Ansible on Linode, Docker, Zitadel for authentication, observability, alerting and incident runbooks. Before joining I worked as a research intern on deepfake detection and as a cybersecurity consulting intern at Ernst & Young India. Three peer-reviewed papers published (Springer ICTIS 2025 and two on IEEE Xplore), two under review. I am seeking an engineering position in Japan and would require visa sponsorship.

活かせる経験・知識・技術 Applicable experience and technical skills

AI・LLM システム / AI and LLM	Agentic systems, MCP orchestration, multi-agent workflows, RAG, LangChain, LangGraph, Claude API, LLM integration, evaluation harnesses
インフラ・DevOps / Infrastructure	CI/CD, OpenTofu (IaC), Ansible, Docker, Linux, Linode, observability and alerting, incident runbooks, Zitadel (SSO/IAM), RBAC, Git
言語・バックエンド / Languages and backend	Python, FastAPI, REST APIs, PostgreSQL, SQL, Swift, Rust, shell scripting
機械学習・画像認識 / ML and vision	PyTorch, TensorFlow, YOLOv8, OpenCV, CNN, ResNet, quantized model deployment

職務経歴 Professional experience — most recent first

Colligence Research	2025年11月 ~ 現在・インド・バンガロール
事業内容	TODO: 会社の事業内容を1行で記入（日本の職務経歴書では会社概要を添えるのが通例）
組織規模・立場	Founding team; four people at the time the deployment work described below was done.
雇用形態	TODO: 雇用形態（正社員／契約社員等）を記入

AI Engineer 2026年2月 ~ 現在

- Work across both halves of the product: the AI and product engineering described below, together with end-to-end ownership of the platform it runs on.
- Established the DevOps/SRE function from scratch and now lead it — CI/CD, infrastructure-as-code with OpenTofu and Ansible on Linode, observability, alerting and incident runbooks.
- Replaced the manual GitHub deploy workflows the team had been using, which reduced deployment errors for the four-person founding team.
- Operate the production stack end to end: Docker, Zitadel authentication, and monitored deployments with self-healing alerting.
- Own system architecture across infrastructure and tooling.
- Building a retrieval-augmented system into the infrastructure layer so that infrastructure-as-code and SRE work — runbook lookup, change context, incident response — is grounded in the organisation's own configuration and history. In progress.

Associate AI Engineer 2025年12月 ~ 2026年2月

- Led a small team building the company's first multi-agent LLM pipeline for ERP automation over MCP, keeping long-running, multi-step workflows on track without human intervention.
- Built a code-generation service that parses structured input specifications and emits deployment-ready web applications.

Associate AI Engineer Intern 2025年11月 ~ 2025年12月

- Prototyped the first LLM integrations and agent workflows, which the production multi-agent pipeline was later built on.

Technologies: Python, FastAPI, MCP, LangChain, LangGraph, Claude API, multi-agent orchestration, OpenTofu, Ansible, Docker, Linode, Linux, Zitadel, CI/CD

Research Intern, AI and Deep Learning

- Built a deepfake-detection model from scratch for manipulated-face video, using no external detection APIs.
- Trained and evaluated the model on the FaceForensics++ and DFDC datasets.

Technologies: Python, PyTorch, TensorFlow, CNN, computer vision, FaceForensics++, DFDC

Ernst & Young India

2024年8月 ~ 2024年9月 ・ インド・バンガロール

Associate Consultant Intern, Enterprise Cybersecurity

- Shadowed senior consultants on cloud-security and identity and access management work.
- Gained exposure to the TOGAF and IFTAS frameworks.

Exposure: Cloud security, IAM, TOGAF, IFTAS

主な担当プロジェクト Selected projects

Arogya-Sathi — offline-capable multilingual AI healthcare platform

2025年1月 ~ 2025年6月

Symptom checking, X-ray fracture detection, prescription summarisation and a wellness chatbot, running fully offline on quantized models and demonstrated on a Raspberry Pi. A deterministic, rule-based drug-interaction and crisis-detection safety layer sits underneath the model and the LLM cannot override it. A 62-case labelled evaluation harness scores those safety layers and runs as a CI merge gate at recall 1.00; building it surfaced a real missed self-harm disclosure. A LangGraph pipeline turns an uploaded medical report into a plain-language explanation in the patient's own language with a self-verification loop, and requests are routed across local Ollama models. Third place, college final-year showcase. github.com/Developer1010x/Arogya-Sathi

Measured: crisis detection recall 1.00 / precision 0.93 ・ emergency 1.00 / 1.00 ・ drug screening 1.00 / 1.00

Technologies: LangGraph, Ollama, quantized models, evaluation harness, CI gate, Raspberry Pi

agentic-devops — AI-powered CI/CD pipeline generator and monitor

Generates and monitors CI/CD pipelines for GitHub and GitLab repositories. github.com/Developer1010x/agentic-devops

Knotes Central — student resource platform

A student resource platform, adopted across 12 or more departments. github.com/Developer1010x/KnotesCentral-Source-Code

pybridge.ai — AI access from a phone

Routes AI access across WhatsApp, Telegram and email so it can be driven from a phone. github.com/Developer1010x/pybridge.ai

LLM Terminal — terminal-based LLM assistant

A lightweight terminal-based assistant for working with LLMs from the command line. github.com/Developer1010x/LLM_Terminal

Audio Deepfake Detection — synthetic speech detection

Flags AI-synthesised or cloned speech. github.com/Developer1010x/Deepfake-Audio-and-AI-Content-Detection

研究業績 Publications

EduConnect: A Smart Solution for Streamlined School Administrations — Springer ICTIS 2025, Bangkok (ICT for Intelligent Systems, Vol. 9).

An Empirical Study of ResNet50 Hyperparameter Tuning for Plant Disease Classification — IEEE Xplore, November 2024. ieeexplore.ieee.org/document/10816835

Web-Server Controlled Rover with Robotic Arm and Object Detection — IEEE Xplore, November 2024. ieeexplore.ieee.org/document/10816816

Two further papers are under review.

学歴 Education

2021年12月 ~ 2025年8月	R V College of Engineering (Visvesvaraya Technological University), Bengaluru, India — B.E. in Computer Science and Engineering
--------------------	---

日本語 / Japanese	TODO: 日本語能力を本人確認のうえ記入（例：学習中／JLPT NO／なし）。確認できない限り記載しない。 Japanese ability is screened on for nearly every role in Japan, including English-speaking ones. Do not enter a level that has not been awarded.
英語 / English	TODO: 英語力の表記を本人と確認（業務使用歴に基づく表現にとどめ、未受験の試験スコアは記載しない）。
資格 / Certifications	TODO: 保有資格。無い場合は「特になし」と記入。

就労資格 Work eligibility in Japan

Currently resident in Bengaluru, India, with no existing right to work in Japan. The applicable status of residence is **Engineer / Specialist in Humanities / International Services** (技術・人文知識・国際業務), the standard category for software engineers hired from overseas; a B.E. in Computer Science is a qualifying degree for it. The Certificate of Eligibility is applied for by the sponsoring employer. Available to relocate to Japan.

自己PR Self-introduction

What I would bring is the combination rather than either half of it: I build the agent and LLM systems and I also run the platform they are deployed on, so the person who ships a workflow is the same person who is responsible for it in production. At Colligence Research that has meant setting up the delivery and monitoring practices a founding team did not yet have, while continuing to deliver product features alongside them. I try to make correctness checkable rather than assumed: in Arogya-Sathi the safety behaviour is enforced by a deterministic layer the model cannot override and verified by a labelled evaluation set that blocks a merge if it regresses. I still have a great deal to learn, and I would welcome the opportunity to do so within your team.

以 上