

S Prajwall Narayana

AI Engineer

Liverpool · eligible for Skilled Worker sponsorship | prajwallnarayana@gmail.com | +91 7760604439 | github.com/Developer1010x | [portfolio](#)

PERSONAL STATEMENT

AI Engineer on the founding team at Colligence Research, working across both halves of the same system. I build the product: multi-agent and LLM pipelines over MCP, a code-generation service, and the backend APIs behind them. I also own the platform it runs on, having established the DevOps and SRE function from scratch and led it since: CI/CD, infrastructure-as-code, Docker, observability, alerting, incident runbooks and authentication. Published three times in peer-reviewed venues (Springer and IEEE). I take ambiguous problems from prototype to a deployed, monitored, instrumented system, and I stay responsible for it afterwards.

Right to work: requires Skilled Worker sponsorship, and qualifies for the **New Entrant** salary rate. Currently in Bengaluru, India; relocating to the UK.

TECHNICAL SKILLS

AI and LLM Systems: Agentic systems, MCP orchestration, multi-agent workflows, RAG, LangChain, LangGraph, Claude API, LLM integration, evaluation harnesses

Infrastructure and DevOps: CI/CD, OpenTofu (IaC), Ansible, Docker, Linux, Linode, observability and alerting, incident runbooks, Zitadel (SSO/IAM), RBAC, Git

Languages and Backend: Python, FastAPI, REST APIs, PostgreSQL, SQL, Swift, Rust, shell scripting

Machine Learning and Vision: PyTorch, TensorFlow, YOLOv8, OpenCV, CNNs, ResNet, quantised model deployment

EXPERIENCE

Colligence Research

Nov 2025 – Present · Bengaluru, India

AI Engineer · Founding Team

Feb 2026 – Present

- Work across both halves of the product: the AI and application engineering below, plus end-to-end ownership of the platform it runs on.
- Established the DevOps and SRE function from scratch and now lead it: CI/CD, infrastructure-as-code with OpenTofu and Ansible on Linode, observability, alerting and incident runbooks.
- Replaced manual GitHub deploy workflows with automated pipelines, cutting deployment errors for a four-person founding team.
- Operate the production stack end to end: Docker, Zitadel authentication and RBAC, and monitored deployments with self-healing alerting.
- Own system architecture across the infrastructure and tooling workstreams.
- Building a retrieval-augmented system into the infrastructure layer so that infrastructure-as-code and SRE work (runbook lookup, change context, incident response) is grounded in the organisation's own configuration and history. In progress.

Associate AI Engineer · Founding Team

Dec 2025 – Feb 2026

- Led a small team building the company's first multi-agent LLM pipeline for ERP automation over MCP, keeping long-running, multi-step workflows on track without human intervention.
- Designed and built a code-generation service that parses structured input specifications and emits deployment-ready web applications, removing manual scaffolding from the loop.

Associate AI Engineer Intern

Nov 2025 – Dec 2025

- Prototyped the founding team's first LLM integrations and agent workflows, the foundation the production multi-agent pipeline was later built on.

Python, MCP, LangChain, LangGraph, Claude API, FastAPI, multi-agent orchestration, OpenTofu, Ansible, Docker, Linode, CI/CD, Zitadel, Linux

RVCE Centre of Excellence

Oct 2024 – Dec 2024 · Remote

Research Intern · AI and Deep Learning

- Built a deepfake-detection model from scratch for manipulated-face video, using no external detection APIs.
- Trained and evaluated the model on the FaceForensics++ and DFDC benchmark datasets.

Python, PyTorch, TensorFlow, CNNs, computer vision, FaceForensics++, DFDC

Ernst & Young India

Aug 2024 – Sep 2024 · Bengaluru, India

Associate Consultant Intern · Enterprise Cybersecurity

- Shadowed senior consultants on enterprise cloud-security and identity and access management engagements.
- Gained exposure to the TOGAF and IFTAS frameworks and to how cloud security reviews are run.

Cloud security, identity and access management, TOGAF, IFTAS

SELECTED ENGINEERING PROJECTS

Arogya-Sathi – *offline-capable multilingual AI healthcare platform*

Jan 2025 – Jun 2025

- Built a deterministic, rule-based drug-interaction and crisis-detection safety layer that the language model cannot override, so clinical safety does not depend on model behaviour.
- Wrote a 62-case labelled evaluation harness that scores those safety layers and wired it in as a continuous-integration merge gate at recall 1.00; building the harness surfaced a real missed self-harm disclosure that the model had let through.
- Built a LangGraph pipeline that turns an uploaded medical report into a plain-language explanation in the patient's own language, with a self-verification loop, and routes requests across multiple local Ollama models.
- Runs fully offline on quantised models: symptom checking, X-ray fracture detection, prescription summarisation and a wellness chatbot, demonstrated end to end on a Raspberry Pi. Placed third at the college final-year showcase.

Measured on the evaluation harness: crisis detection recall 1.00 / precision 0.93 · emergency 1.00 / 1.00 · drug screening 1.00 / 1.00LangGraph, Ollama, quantised models, evaluation harness, CI merge gate, Raspberry Pi · github.com/Developer1010x/Arogya-Sathi**agentic-devops** – *AI-powered CI/CD pipeline generator and monitor*

- Generates and monitors CI/CD pipelines for GitHub and GitLab repositories, applying the agent work to the delivery pipeline itself.

CI/CD, GitHub, GitLab, agents · github.com/Developer1010x/agentic-devops**pybridge.ai** – *AI access routed to a phone*

- Routes AI access across WhatsApp, Telegram and email so the whole system can be driven from a phone.

Python, messaging, LLM integration · github.com/Developer1010x/pybridge.ai**Knotes Central** – *student resource platform*

- Co-built a platform delivering centralised notes and academic material, adopted across 12+ departments.

Web platform · github.com/Developer1010x/KnotesCentral-Source-Code**Audio Deepfake Detection** – *synthetic speech detection*

- Model that flags AI-synthesised or cloned speech.

PyTorch, audio, classification · github.com/Developer1010x/Deepfake-Audio-and-AI-Content-Detection**LLM Terminal** – *terminal-based LLM assistant*

- Lightweight assistant for working with language models from the command line.

Python, CLI, LLM integration · github.com/Developer1010x/LLM_Terminal

PUBLICATIONS

EduConnect: A Smart Solution for Streamlined School Administrations – Springer ICTIS 2025, Bangkok (ICT for Intelligent Systems, Vol. link.springer.com/book/10.1007/978-981-96-8750-3 9).**An Empirical Study of ResNet50 Hyperparameter Tuning for Plant Disease Classification** – IEEE Xplore, Nov 2024.ieeexplore.ieee.org/document/10816835**Web-Server Controlled Rover with Robotic Arm and Object Detection** – IEEE Xplore, Nov 2024.ieeexplore.ieee.org/document/10816816 Two further papers are under review.

EDUCATION

R V College of Engineering (VTU)

Dec 2021 – Aug 2025 · Bengaluru, India

B.E. Computer Science and Engineering