

S Prajwall Narayana

AI Engineer

Bengaluru, India (open to US relocation) | prajwallnarayana@gmail.com | +91 7760604439 |
github.com/Developer1010x | sprajwallnarayana.vercel.app

SUMMARY

AI Engineer on a startup founding team, working across both halves of the product: the agentic and LLM systems that ship to users, and the CI/CD, infrastructure-as-code, observability and authentication platform they run on. Designed a deterministic safety layer under a healthcare LLM and a 62-case labeled evaluation harness that gates merges in CI at recall 1.00. Three peer-reviewed papers (IEEE, Springer). **Requires US visa sponsorship (H-1B, or F-1/OPT if studying in the US).**

TECHNICAL SKILLS

AI and LLM Systems: Agentic Systems, Multi-Agent Workflows, MCP Orchestration, LangGraph, LangChain, Claude API, RAG, Evaluation Harnesses, LLM Integration

Infrastructure and DevOps: CI/CD, OpenTofu (IaC), Ansible, Docker, Observability and Alerting, Incident Runbooks, Zitadel (SSO/IAM), RBAC, Linux, Linode, Git

Languages and Backend: Python, FastAPI, REST APIs, PostgreSQL, SQL, Shell Scripting, Rust, Swift

Machine Learning and Vision: PyTorch, TensorFlow, Quantized Model Deployment, YOLOv8, OpenCV, CNN, ResNet

EXPERIENCE

Colligence Research – Bengaluru, India

AI Engineer, Founding Team | Feb 2026 – Present

- Own both halves of the product: the AI and LLM systems that ship to users, and the platform they run on.
- Cut deployment errors for a four-person founding team by replacing manual GitHub deploy workflows with automated CI/CD; established the DevOps/SRE function from scratch and now lead it.
- Run the production stack end to end (Docker, Zitadel authentication, monitored deployments with self-healing alerting) on infrastructure defined as code with OpenTofu and Ansible on Linode, with observability, alerting and incident runbooks; own system architecture across infrastructure and tooling.
- Building a retrieval-augmented system into the infrastructure layer so that infrastructure-as-code and SRE work (runbook lookup, change context, incident response) is grounded in the organization's own configuration and history. In progress.

Associate AI Engineer, Founding Team | Dec 2025 – Feb 2026

- Led a small team building the company's first multi-agent LLM pipeline for ERP automation over MCP, keeping long-running, multi-step workflows on track without human intervention.
- Built a code-generation service that parses structured input specs and emits deployment-ready web applications.

Associate AI Engineer Intern | Nov 2025 – Dec 2025

- Prototyped the company's first LLM integrations and agent workflows.

Tools: Agentic Systems, MCP, LangGraph, FastAPI, Python, OpenTofu, Ansible, Docker, Linode, Zitadel, CI/CD

RVCE Centre of Excellence – Remote

Research Intern, AI and Deep Learning | Oct 2024 – Dec 2024

- Built a deepfake-detection model from scratch for manipulated-face video, using no external detection APIs; trained and evaluated it on the FaceForensics++ and DFDC datasets.

Tools: PyTorch, TensorFlow, CNN, FaceForensics++, DFDC

Ernst & Young India – Bengaluru, India

Associate Consultant Intern, Enterprise Cybersecurity | Aug 2024 – Sep 2024

- Shadowed senior consultants on cloud-security and identity and access management work; exposure to the TOGAF and IFTAS frameworks.

Exposure: Cloud Security, IAM, TOGAF, IFTAS

PROJECTS

Arogya-Sathi, Offline-Capable Multilingual AI Healthcare Platform

Jan 2025 – Jun 2025 | github.com/Developer1010x/Arogya-Sathi

- Built a 62-case labeled evaluation harness scoring the platform's safety layers and wired it into CI as a merge gate at recall 1.00: crisis detection 1.00 recall / 0.93 precision, emergency 1.00 / 1.00, drug screening 1.00 / 1.00. Building the harness surfaced a real missed self-harm disclosure.
- Designed the deterministic, rule-based drug-interaction and crisis-detection safety layer that the harness scores, sitting underneath the model so the LLM cannot override it.
- Shipped a LangGraph pipeline that turns an uploaded medical report into a plain-language explanation in the patient's own language with a self-verification loop, routing requests across multiple local Ollama models.
- Product surface: symptom checking, X-ray fracture detection, prescription summarization and a wellness chatbot, running fully offline on quantized models and demonstrated on a Raspberry Pi. Third place, college final-year showcase.

Tools: LangGraph, Ollama, Quantized Models, Evaluation Harness, CI Merge Gate, Raspberry Pi

Open-Source and Personal Systems

github.com/Developer1010x

- agentic-devops: AI-powered CI/CD pipeline generator and monitor for GitHub and GitLab.
- pybridge.ai: routes AI access across WhatsApp, Telegram and email so it can be driven from a phone.
- LLM Terminal: lightweight terminal-based LLM assistant. Audio Deepfake Detection: flags AI-synthesized or cloned speech.
- Knotes Central: student resource platform, adopted across 12+ departments.

PUBLICATIONS

EduConnect: A Smart Solution for Streamlined School Administrations – Springer ICTIS 2025, Bangkok (ICT for Intelligent Systems, Vol. link.springer.com/book/10.1007/978-981-96-8750-3 9).

An Empirical Study of ResNet50 Hyperparameter Tuning for Plant Disease Classification – IEEE Xplore, Nov 2024. ieeexplore.ieee.org/document/10816835

Web-Server Controlled Rover with Robotic Arm and Object Detection – IEEE Xplore, Nov 2024. ieeexplore.ieee.org/document/10816816. Two more papers under review.

EDUCATION

B.E. Computer Science and Engineering, R V College of Engineering (VTU), Bengaluru, India | Dec 2021 – Aug 2025.